

Matchmaker app on Google Cloud

השדכן החכם שלי


A person uploads a CV and gets the open Israeli Civil Service jobs that fit it, with a score and a short explanation. The website is static. One container runs the API. Gemini is the existing AI Studio API.

Services required

Service	What it is for
Firebase Hosting	The website. Serves the built Vue app on the Firebase CDN.
Cloud Run	The API. One container, service name <code>matchmaker-api</code> , region <code>me-west1</code> (Tel Aviv).
Artifact Registry	Stores the API container image that Cloud Run runs. Same region.
Secret Manager	Stores the Gemini API key and passes it to Cloud Run as <code>GOOGLE_GENERATIVE_AI_API_KEY</code> .
Billing	A billing account on the Firebase project (Blaze). Cloud Run cannot run without it.

Not created in this project: **Gemini API** (Google AI Studio, same key and models as today, not Vertex AI) and the public **Merkava** jobs API. There is no database and no file store.

Request flow



How a visit runs

Firebase Hosting only serves files. The browser then calls Cloud Run directly at the URL baked into the build (VITE_API_URL). Cloud Run allows that Hosting origin with CORS.

Do not send /api through a Hosting rewrite. Hosting stops those requests at 60 seconds, and a match takes longer. Cloud Run is set to 300 seconds and streams progress back while Gemini works.

On a match, the API reads the CV, asks Gemini for a profile, screens every open job, then scores the best matches. The CV is held in memory for that request only. It is not saved.

Firebase Hosting

- Publish the built web app (apps/web/dist).
- Cache hashed files under /assets for one year. Send every other path to index.html.
- Build with VITE_API_URL set to https://<cloud-run-host>/api.
- Default hosts are *.web.app and *.firebaseapp.com. A custom domain can be added on Hosting later.

Cloud Run · matchmaker-api

Setting	Value
Region	me-west1 (Tel Aviv), close to users and to Merkava
Access	Public HTTPS. Grant allUsers the invoker role.
Port	8080. The process already listens on PORT.
Timeout	300 seconds
Size	1 vCPU, 1 GiB, concurrency 10
Instances	Minimum 1 (keeps the job list warm), maximum 3
Health	GET /api/health
Routes	/api/health, /api/positions, /api/positions/:id, POST /api/match

Secret Manager

- One secret: the existing AI Studio key.
- Cloud Run reads it as GOOGLE_GENERATIVE_AI_API_KEY.
- The runtime service account needs Secret Accessor on that secret.
- The key is not in the image and not in the website.

Artifact Registry

- One Docker repository in me-west1.
- Cloud Run deploys the image from there.
- The image is Node 22 and the Hono API only. It does not include the website.

API environment

Variable	Purpose
GOOGLE_GENERATIVE_AI_API_KEY	From Secret Manager. Required.
ALLOWED_ORIGINS	The Hosting URLs, so the browser may call the API.
GEMINI_MODEL	Optional. Default gemini-3.5-flash-lite, used to read the CV and screen jobs.
GEMINI_SCORING_MODEL	Optional. Defaults to the same model. Scoring step only.